



WEEK-VIII

TYPES OF MODELING



CLASSIFICATION (Supervised Learning)

- ② structural assumptions
- GENERATIVE MODEL
 - DISCRIMINATIVE MODEL
- modelling



Q. Which one is better?

* Knowledge of how the data is generated helps us discriminate between the classes.

$$\text{DATA} = \{ (x_1, y_1) \dots (x_n, y_n) \}$$

$x_i \in \{0, 1\}^d \longrightarrow \text{Binary features}$
 $y_i \in \{0, 1\}$

EXAMPLE: SPAM CLASSIFICATION

$$y = \{0, 1\}$$

↑
Spam
↓
Non-spam

$$x_i \in \{0, 1\}^d$$

words in dictionary

Eg: "Hello, how are you?"

$$\rightarrow \begin{bmatrix} \text{About} & & \text{ARE} & & \text{HELLO} & & \text{How} & & \text{you} & \text{ZEBRA} \\ 0 & 0 & 0 & \dots & 1 & 1 & 1 & \dots & 1 & 0 \end{bmatrix} \in \mathbb{R}^d$$

$$P(x, y) \quad ("Hello, how are you?", \text{spam}) = 0.01$$

modelling the structure of an email is not necessary; one can choose a Discriminative model to model $p(\text{spam/email})$

$$P(x, y) = \underbrace{P(x)}_{p(\text{email})} \cdot \underbrace{P(y/x)}_{p(\text{spam/email})}$$

$$= \underbrace{P(y)}_{p(\text{spam})} \cdot \underbrace{P(x/y)}_{p(\text{email/spam})}$$

we care about $p(\text{spam/email})$ only to make prediction.

likelihood of seeing a spam or a non-spam

structure of a particular email given spam

GENERATIVE STORY

STEP-1: DECIDE the Label by tossing a coin

$$P(y_i = 1) = p$$

$$P(y_i = 0) = (1 - p)$$

STEP 2:

DECIDE FEATURES using the label in step 1

$$\text{by } P(x_i / y_i)$$

* features are not independent of each other.


d=3

events	w ₁	w ₂	w ₃
0 0 0	0.1		
0 0 1	0.01		
0 1 0	0.02		
0 1 1	0.03		
1 0 0	0.04		
1 0 1	0.5		
1 1 0	0.2		
1 1 1	0.1		

$\sum p = 1$

don't need this element's probability as the total probability is 1.

$2^3 = 8$ different feature vectors as $d=3$



events	w ₁	w ₂	w ₃
0 0 0	0	0	0
0 0 1	0	0	1
0 1 0	0	1	0
0 1 1	0	1	1
1 0 0	1	0	0
1 0 1	1	0	1
1 1 0	1	1	0
1 1 1	1	1	1

$\sum p = 1$

$2^3 = 8$ different feature vectors as $d=3$

parameters in the model = 1 + $\underbrace{(2^d - 1)}_{\substack{\text{to decide label} \\ P(x|y=1)}} + \underbrace{(2^d - 1)}_{\substack{\text{free parameters} \\ P(x|y=0)}}$

= $1 + 2(2^d - 1)$

= $2^{d+1} - 1$

• Multiclass Classification Problem:

K-classes - $y_i \in \{1, 2, 3, \dots, K\}$

d-binary features - $x_i \in \{0, 1\}^d$

parameters to be estimated = parameters in the model

= $(K-1) + K(2^d - 1)$

= $K(2^d) - 1$

ISSUE

- Too many parameters!
- Not a reasonable story.


d=3

SPAM

0 0 0
0 0 1
0 1 0
0 1 1
1 0 0
1 0 1
1 1 0
1 1 1

TAIL HEAD HEAD

d-coins:



Non-spam

0 0 0
0 0 1
...
1 1 1

2^d

: d-coins

★ CLASS CONDITIONAL INDEPENDENCE ASSUMPTION!

parameters = $1 + d + d = 2d + 1$

MANAGEABLE!

parameters to be estimated = parameters in the model

= $(K-1) + Kd$

Example →

$$\begin{aligned}
 & y=1 \\
 & \begin{array}{c} \textcircled{p_1'} \quad \textcircled{p_2'} \quad \textcircled{p_3'} \\ \swarrow \quad \downarrow \quad \searrow \\ p(x=[0 \ 1 \ 0] / y=1) \end{array} \\
 & = \\
 & (1-p_1') \cdot (p_2') \cdot (1-p_3')
 \end{aligned}$$

STEP 1:

$$P(y=1) = p$$

$$P(y=0) = 1-p$$

STEP 2:

$$P(x=[f_1 \ f_2 \ \dots \ f_d] / y) ; \begin{matrix} f_i \in \{0,1\} \\ y_i \in \{0,1\} \end{matrix}$$

$$= \prod_{i=1}^d \left(p_i^y (1-p_i^y)^{(1-f_i)} \right)$$

iid
FEATURES ARE
CONDITIONALLY INDEPENDENT
GIVEN LABEL


• How to estimate the parameters

• SOLUTION: MAXIMUM LIKELIHOOD!

- ★ The primary goal of a generative model is to find the parameters that maximize the probability (likelihood) of the training data occurring.

Calculating joint probabilities requires multiplying many tiny fractions, which causes computational underflow in computers. Taking the logarithm solves this by converting multiplication into addition, keeping the math numerically stable without changing the final optimal result.

Standard machine learning algorithms are built to minimise an Objective (cost/loss) function, not maximise it. By taking the negative of the log-likelihood, we perfectly flip our maximisation goal into a standard minimisation problem.



CLASS CONDITIONAL INDEPENDENCE

$2d+1$
parameters in the model

PARAMETER ESTIMATION

$$p, \underbrace{\{p'_1, \dots, p'_d\}}_{\text{feature}}, \underbrace{\{p''_1, \dots, p''_d\}}_d$$

label

MAX. LIKELIHOOD ESTIMATES

$$\hat{p} = \frac{1}{n} \sum_{i=1}^n y_i$$

$1 \rightarrow$
parameter

sample mean of the labels

$\{y=1\}$ -labelled datapoints
the dataset

$\hat{p}$ is trying to estimate the probability whether an email is spam or not.

$2d \rightarrow$
parameters

$$\hat{p}_j^y = \frac{\sum_{i=1}^n \mathbb{1}(f_j^i = 1, y_i = y)}{\sum_{i=1}^n \mathbb{1}(y_i = y)}$$

$\forall j \in \{1, \dots, d\}$
 $\forall y \in \{0, 1\}$
 $\forall i \in \{1, \dots, n\}$

$$\frac{\sum_{i=1}^n \mathbb{1}(f_j^i = 1, y_i = y)}{\sum_{i=1}^n \mathbb{1}(y_i = y)}$$

j^{th} feature in the email
 i^{th} datapoint

Number of datapoints with label y .

Fraction of y -labelled datapoints that contain the j^{th} word.

PREDICTION

Given $x^{\text{test}} \in \{0, 1\}^d$, what is $\hat{y}^{\text{test}}$?

$$P(y^{\text{test}} = 1 | x^{\text{test}}) > P(y^{\text{test}} = 0 | x^{\text{test}})$$

$$\Rightarrow \hat{y}^{\text{test}} = 1$$

$= 0$ otherwise.

How to obtain $P(y/z)$ from $P(y)$ and $P(z/y)$?

BAYES' RULE! $\rightarrow P(A|B) = \frac{P(B|A)P(A)}{P(B)}$

$$P(y^{\text{test}} = 1 | x^{\text{test}}) = \frac{P(x^{\text{test}} | y^{\text{test}} = 1) \cdot P(y^{\text{test}} = 1)}{P(x^{\text{test}})}$$

$P(x^{\text{test}})$ $\rightarrow$ +ve denominator does not affect comparison.

$$P(y^{\text{test}} = 0 | x^{\text{test}}) = \frac{P(x^{\text{test}} | y^{\text{test}} = 0) \cdot P(y^{\text{test}} = 0)}{P(x^{\text{test}})}$$

$P(x^{\text{test}})$ $\rightarrow$ +ve denominator does not affect comparison.

$$P(x^{\text{test}} | y^{\text{test}} = 1) \cdot P(y^{\text{test}} = 1)$$

$$= P(x^{\text{test}} = [f_1, f_2, \dots, f_d] | y^{\text{test}} = 1) \cdot P(y^{\text{test}} = 1)$$

$$= \left(\prod_{j=1}^d (\hat{p}_j^{f_j} (1 - \hat{p}_j)^{(1-f_j)}) \right) \cdot \hat{p}$$

d-parameters + 1-parameter

• features are independent of each other

using MLE here $\rightarrow$

$$x_{\text{test}} = \begin{bmatrix} f_1 & f_2 & f_3 & f_4 \\ 0 & 0 & 1 & 0 \end{bmatrix}$$

$$\left(\prod_{j=1}^d (\hat{p}_j^{f_j} (1 - \hat{p}_j)^{(1-f_j)}) \right) \cdot \hat{p} > \left(\prod_{j=1}^d (\hat{p}_j^{f_j} (1 - \hat{p}_j)^{(1-f_j)}) \right) \cdot (1 - \hat{p})$$

$\Rightarrow$ PREDICT $\hat{y}^{\text{test}} = 1$
ELSE $\hat{y}^{\text{test}} = 0$

NAIVE BAYES ALGORITHM.

MODEL USES 2 main assumptions Things $\rightarrow$ features are independent given the labels

- $2^d \rightarrow 2d$
- CLASS CONDITIONAL INDEPENDENCE (iid)

• BAYES' THEOREM

- may not hold in practice
- NAIVE assumption
- Still works well in practice, especially when d is large.

PITFALLS IN NAIVE BAYES TO WATCH OUT FOR.

- If a word does not appear in the train set but appears in a test datapoint,

$$\hat{p}_i^1 = 0 \quad \hat{p}_i^0 = 0$$

$$P(y^{\text{test}} = 1 \mid x^{\text{test}} = [f_1, f_2, \dots, f_d]) \propto \left(\prod_{i=1}^d \underbrace{\hat{p}_i^1}_{\substack{\uparrow \\ \text{Zebra} \\ = 1}} \underbrace{(1 - \hat{p}_i^1)}_1 \right) \hat{p} \quad \underbrace{\hspace{10em}}_0$$

$$P(y^{\text{test}} = 0 \mid x^{\text{test}} = [f_1, \dots, f_d]) \propto \left(\prod_{i=1}^d \underbrace{\hat{p}_i^0}_{\substack{\uparrow \\ \text{Zebra} \\ = 1}} \underbrace{(1 - \hat{p}_i^0)}_1 \right) (1 - \hat{p}) \quad \underbrace{\hspace{10em}}_0$$

$$P(y^{\text{test}} = 0 \mid x^{\text{test}}) = P(y^{\text{test}} = 1 \mid x^{\text{test}}) = 0$$

Possible Fix

- Can add two "pseudo" datapoints with all words present - one email has label 0 and another has label 1

LAPLACE SMOOTHING

$$\begin{array}{c} x_1^{\text{pseudo}} \\ [1 \quad 1 \quad 1 \quad 1 \quad \dots \quad 1] \\ x_0^{\text{pseudo}} \\ [1 \quad 1 \quad 1 \quad 1 \quad \dots \quad 1] \end{array} \quad \begin{array}{c} y_1^{\text{pseudo}} \\ 1 \\ 0 \end{array}$$

Every feature appears at least once in both the classes

Probability estimates from maximum likelihood cannot be zero.

DECISION FUNCTION OF NAIVE BAYES

$$\text{Given } x_{\text{test}}; \quad y_{\text{test}} = 1 \quad \text{if} \quad \frac{P(y_{\text{test}} = 1 \mid x_{\text{test}})}{P(y_{\text{test}} = 0 \mid x_{\text{test}})} \geq 1$$

$$\log \left(\frac{P(y_{\text{test}} = 1 | x_{\text{test}})}{P(y_{\text{test}} = 0 | x_{\text{test}})} \right) \geq 0 \quad (\log(1) = 0)$$

$$\log \left(\frac{P(x_{\text{test}} | y_{\text{test}}=1) \cdot P(y_{\text{test}}=1)}{P(x_{\text{test}} | y_{\text{test}}=0) \cdot P(y_{\text{test}}=0)} \right) \geq 0$$

$$X_{\text{test}} = [f_1 \ f_2 \ \dots \ f_d]$$

$$\log \left(\frac{\prod_{i=1}^d \frac{(\hat{p}_i')^{f_i} (1 - \hat{p}_i')^{(1-f_i)}}{(\hat{p}_i^0)^{f_i} (1 - \hat{p}_i^0)^{(1-f_i)}} \cdot \hat{p}}{(1-\hat{p})} \right) \geq 0$$

$$\therefore \log \left(\prod_{i=1}^d \left(\frac{\hat{p}_i^1}{\hat{p}_i^0} \right)^{f_i} \left(\frac{1 - \hat{p}_i^1}{1 - \hat{p}_i^0} \right)^{(1-f_i)} \cdot \frac{\hat{p}}{1 - \hat{p}} \right) \geq 0$$

$$\sum_{i=1}^d \left(f_i \log \left(\frac{\hat{p}_i^1}{\hat{p}_i^0} \right) + (1-f_i) \log \left(\frac{1-\hat{p}_i^1}{1-\hat{p}_i^0} \right) + \log \left(\frac{\hat{p}_i^1}{1-\hat{p}_i^1} \right) \right) \geq 0$$

$$= \sum_{i=1}^d \underbrace{f_i \left(\log \left(\frac{\hat{p}_i' (1 - \hat{p}_i^0)}{\hat{p}_i^0 (1 - \hat{p}_i') } \right) \right)}_{\text{red}} + \underbrace{\log \left(\frac{(1 - \hat{p}_i')}{(1 - \hat{p}_i^0)} \right) + \log \left(\frac{\hat{p}_i^0}{1 - \hat{p}_i^0} \right)}_{\text{blue}} \geq 0$$

$$\log \frac{m}{n} = \log m - \log n$$

$$x_{\text{test}} = [f_1 \dots f_d]$$

DECISION FUNCTION is of the form


Predict $y_{\text{test}} = 1$ if $W^T x_{\text{test}} + b \geq 0$
 $\uparrow$
 $\in \mathbb{R}^d$

Where $w_i = \log \left(\frac{\hat{p}_i^1 (1 - \hat{p}_i^0)}{\hat{p}_i^0 (1 - \hat{p}_i^1)} \right)$, $b =$

$$\{x: P(Y=0|x) = P(Y=1|x)\}$$

CONCLUSION:

→ DECISION FUNCTION OF NAIVE BAYES IS LINEAR!



↓
Bias: how much away from the origin

$$\text{DATA: } \{ (x_1, y_1), \dots, (x_n, y_n) \}$$

$$x_i \in \mathbb{R}^d \quad y_i \in \{0, 1\}$$

A GENERATIVE STORY

Σ tells us how the density contours look like.
regions w/ same density


Σ tells us how features are co-related.

In a multivariate Gaussian distribution, the shape of the density contours is controlled by the inverse covariance matrix, also called the Precision matrix.

$$\Sigma^{-1}$$

PARAMETERS

- p ellipses around the mean and the direction of ellipses is determined by Σ .
- μ_0, μ_1
- Σ : co-variance matrix, common for both classes



NOTE: In this model, covariances are assumed to be same

MAXIMUM LIKELIHOOD ESTIMATES

$$\hat{p} = \frac{\sum_{i=1}^n y_i}{n}$$

← FRACTION OF points labelled 1.

$$\hat{\mu}_1 = \frac{\sum_{i=1}^n \mathbb{1}(y_i = 1) \cdot x_i}{\sum_{i=1}^n \mathbb{1}(y_i = 1)}$$

← Sample mean of data points labelled 1.

$$\hat{\mu}_0 = \frac{\sum_{i=1}^n \mathbb{1}(y_i = 0) \cdot x_i}{\sum_{i=1}^n \mathbb{1}(y_i = 0)}$$

← Sample mean of data points labelled 0.

$$\hat{\Sigma} = \frac{1}{n} \sum_{i=1}^n (x_i - \hat{\mu}_{y_i})(x_i - \hat{\mu}_{y_i})^T$$

y_i is the label associated with x_i

PREDICTION? Bayes rule.

$$P(y_{\text{test}} | x_{\text{test}}) \propto \underbrace{P(x_{\text{test}} | y_{\text{test}})}_{f(x_{\text{test}}; \hat{\mu}_{y_{\text{test}}}, \hat{\Sigma})} \cdot \underbrace{P(y_{\text{test}})}_{\hat{p}}$$

Predict $y_{\text{test}} = 1$ if

$$f(x_{\text{test}}; \hat{\mu}_1, \hat{\Sigma}) \cdot \hat{p} \geq f(x_{\text{test}}; \hat{\mu}_0, \hat{\Sigma}) \cdot (1 - \hat{p})$$

$$\frac{-(x_{\text{test}} - \hat{\mu}_1)^T \hat{\Sigma}^{-1} (x_{\text{test}} - \hat{\mu}_1)}{e^{\hat{p}}} \geq \frac{-(x_{\text{test}} - \hat{\mu}_0)^T \hat{\Sigma}^{-1} (x_{\text{test}} - \hat{\mu}_0)}{e^{(1-\hat{p})}}$$

on simplification


[take log]

Predict 1 if

$$\underbrace{\left(\hat{\mu}_1 - \hat{\mu}_0 \right)^T \hat{\Sigma}^{-1}}_{w^T} x_{\text{test}} + b \geq 0$$



not dependant on data

$$\hat{\mu}_0^T \hat{\Sigma}^{-1} \hat{\mu}_0 - \hat{\mu}_1^T \hat{\Sigma}^{-1} \hat{\mu}_1 + \log\left(\frac{1-\hat{p}}{\hat{p}}\right) \geq 0$$



DECISION FUNCTION IS LINEAR!

$\rightarrow$ $\star \Sigma$ is same for both classes.



GAUSSIAN NAIVE BAYES